

ibf-thresholds-triggers

Release v1

Nishadh Kalladath, Eunice Koech, Anthony Mwanthi

Apr 06, 2025

1	Anticipatory Action for Drought Early Warning	3
1.1	Concept of Anticipatory Action (AA)	3
1.2	Ensemble Prediction Systems	3
1.3	Empirical Probability Calculation	3
1.4	Role of Long-term Verification	4
1.5	AUROC: A Key Metric for Forecast Verification	4
1.6	Quality Threshold for Anticipatory Action	4
1.7	Bootstrapping for Uncertainty Assessment	5
1.8	Role in Trigger Selection	5
2	SPI Calculation and Foundational Libraries	7
2.1	The Standardized Precipitation Index (SPI)	7
2.2	The xclim Library: Climate Indices Made Accessible	7
2.3	The xarray Foundation: Multi-dimensional Data Structures	8
3	Regridding Datasets with xESMF	9
3.1	Spatial Resolution Harmonization	9
3.2	The xESMF Library	10
3.3	Handling Special Cases in Regridding	10
4	Aligning Forecast and Observation Data with climpred	11
4.1	Time Alignment Challenge in Forecast Verification	11
4.2	The climpred Library	11
4.3	Lead Time Calculation Example: JJA Season SPI	12
4.4	Seasonal SPI Product Identification	12
5	Forecast Verification with xhistogram	15
5.1	Multi-dimensional Contingency Analysis	15
5.2	Constructing Contingency Tables	15
5.3	Verification Metrics Explained	16
5.4	AUROC and Bootstrapping	17
5.5	2D vs 1D Verification Methods	18
6	Trigger Selection for Anticipatory Action against Droughts in East Africa	19
6.1	Forecast Verification and Trigger Selection Process	19
6.2	Verification Metrics	20
6.3	Downstream Processing Dependencies	20
6.4	Forecast Skill Assessment	21
6.5	Operational Implementation	21
6.6	Time Window Considerations	21

7	Map Visualization for Ensemble Forecasts and Observations	23
7.1	Setting Up the Environment	23
7.2	Data Organization with DataTree.	23
7.3	Creating Single-Row Time Series Plot	24
7.4	Plotting Map Elements	25
7.5	Processing Multiple Initializations	26
7.6	Merging Individual Plots	27
7.7	Example Usage	28
7.8	Conclusion.	29

Add your content using `reStructuredText` syntax. See the [reStructuredText](#) documentation for details.

Anticipatory Action for Drought Early Warning

1.1 Concept of Anticipatory Action (AA)

Anticipatory Action (AA) represents a paradigm shift in humanitarian response, where actions are taken before a disaster occurs based on forecasts rather than after impacts materialize. In drought contexts, AA leverages the slow-onset nature of droughts to implement interventions during the critical “window of opportunity” between a reliable forecast and the onset of drought impacts.

According to Guimarães Nobre et al. (2023), AA enables the implementation of predefined actions triggered by forecasts that exceed specific thresholds, facilitating early risk reduction. This approach links seasonal forecasting information with predetermined contingency plans, actors, and funding—often referred to as Forecast-based Financing (FbF). Unlike traditional humanitarian responses that mobilize after a crisis, AA systems preposition resources and initiate interventions when forecast indicators reach trigger values, potentially reducing losses to livelihoods and infrastructure while providing more dignified support to communities.

1.2 Ensemble Prediction Systems

Ensemble Prediction Systems (EPS) form the backbone of modern seasonal forecasting, addressing inherent uncertainties in weather and climate prediction. These systems generate multiple forecasts (ensemble members) by varying initial conditions, model physics, or both. The European Centre for Medium-Range Weather Forecasts’ SEAS5 system, for instance, provides 25-51 ensemble members with lead times up to 7 months.

As described by Guimarães Nobre et al. (2024), ensemble forecasts allow for probabilistic interpretation of potential outcomes, crucial for risk-based decision-making. Each ensemble member represents an equally likely future scenario, collectively providing information about forecast uncertainty and the range of possible outcomes. This probabilistic approach is essential for AA systems where decisions involve significant resource allocation under uncertainty.

1.3 Empirical Probability Calculation

For drought monitoring and forecasting, empirical probability calculations convert ensemble forecasts into decision-relevant metrics. The process involves:

1. Calculating drought indicators (e.g., Standardized Precipitation Index) for each ensemble member
2. Determining what proportion of ensemble members exceed predefined thresholds
3. Deriving empirical probabilities that represent the likelihood of drought conditions

These empirical probabilities become the basis for triggers in AA systems. According to Guimarães Nobre et al. (2023), the empirical probability approach enables quantitative assessment of forecast uncertainty, allowing decision-makers to balance the risks of action versus inaction.

1.4 Role of Long-term Verification

Long-term verification forms the cornerstone of trust and justification for any AA system. By systematically comparing historical forecasts with observed outcomes over many years, verification:

1. Quantifies forecast quality through metrics like hit rates and false alarm ratios
2. Establishes appropriate trigger values based on performance statistics
3. Builds confidence among stakeholders and donors in the AA approach
4. Supports continuous improvement of the forecasting system

As emphasized in both papers, verification analysis provides the evidence base needed to justify anticipatory investments. It ensures that actions are triggered at appropriate probability thresholds that balance the risks of missed events against false alarms—a critical consideration when resources are limited and each intervention must be carefully justified.

1.5 AUROC: A Key Metric for Forecast Verification

The Area Under the Receiver Operating Characteristic curve (AUROC) represents a fundamental metric in the verification of probabilistic forecasts for Anticipatory Action (AA) systems. This metric addresses a central question: does the forecasting system effectively discriminate between drought and non-drought events?

1.5.1 AUROC Concept and Calculation

The AUROC score quantifies a forecast's ability to discriminate between binary outcomes (drought /no drought) across different probability thresholds. As described by Guimarães Nobre et al. (2023), the calculation involves setting a range of probability thresholds for converting forecasts into binary predictions, then evaluating how well these predictions match observations.

The resulting curve plots the hit rate (proportion of observed droughts correctly forecast) against the false alarm rate (proportion of non-droughts incorrectly forecast as droughts) at various threshold levels. The area under this curve—the AUROC score—ranges from 0 to 1, where:

- 0.5 indicates no skill (equivalent to random guessing)
- 1.0 indicates perfect discrimination
- Values below 0.5 suggest systematically incorrect forecasts

1.6 Quality Threshold for Anticipatory Action

In operational AA systems for drought, an AUROC score of 0.5 serves as a critical quality threshold. As implemented in the “Ready, Set & Go!” system (Guimarães Nobre et al., 2024), forecasts with AUROC scores below 0.5 are considered unsuitable for triggering anticipatory actions, as they provide no added value over random chance.

This threshold represents a minimum standard for forecast quality—one that ensures actions are based on information with at least some predictive skill. When evaluating potential triggers for AA, only those derived from forecasts with $AUROC > 0.5$ are considered, creating a first-level quality control in the system design.

1.7 Bootstrapping for Uncertainty Assessment

To account for sampling uncertainty in AUROC estimation, the verification process employs bootstrapping techniques. By repeatedly resampling the forecast-observation pairs with replacement (typically 1,000 iterations), a distribution of AUROC scores is generated, allowing for:

1. Estimation of confidence intervals around the AUROC score
2. Assessment of statistical significance of skill above the 0.5 threshold
3. More robust decision-making about forecast quality

This approach acknowledges that verification statistics based on limited historical records contain inherent uncertainty, particularly for rare events like severe droughts.

1.8 Role in Trigger Selection

Beyond serving as a quality filter, AUROC scores influence the selection of optimal trigger values. Higher AUROC scores indicate greater potential for skillful discrimination, allowing for more confident selection of triggers based on other performance metrics like hit rates and false alarm ratios.

In the context of long-term verification, AUROC provides a foundation for justifying AA initiatives to stakeholders and funders by demonstrating that the underlying forecasting system possesses genuine predictive skill, making it a suitable basis for humanitarian decision-making.

SPI Calculation and Foundational Libraries

2.1 The Standardized Precipitation Index (SPI)

The Standardized Precipitation Index (SPI) serves as a primary drought indicator in the anticipatory action system, quantifying rainfall anomalies across different timescales. Originally developed by McKee et al. (1993), SPI represents how current precipitation compares to historical conditions, expressed in standard deviation units.

In the implementation found in `run-process-spi.py`, SPI is calculated for 3-month accumulation periods (SPI-3), capturing seasonal-scale rainfall anomalies particularly relevant to agricultural drought:

```
spi_3 = standardized_precipitation_index(  
    aa,  
    freq="MS",  
    window=3,  
    dist="gamma",  
    method="APP",  
    cal_start='1991-01-01',  
    cal_end='2018-01-01',  
    fitkwargs={"floc": 0}  
)
```

SPI values typically range from -4 to +4, with negative values indicating precipitation deficits. The interpretation follows standard thresholds, where values below -1 represent severe drought conditions that occur approximately once every 6-7 years, making them suitable targets for anticipatory action.

2.2 The xclim Library: Climate Indices Made Accessible

The xclim library provides a robust framework for climate data analysis, with particular strengths in calculating standardized climate indices like SPI. Developed to work seamlessly with xarray data structures, xclim implements climate indices following scientific standards while addressing computational challenges associated with large climate datasets.

For drought monitoring, xclim's `standardized_precipitation_index` function offers several critical features:

1. **Distribution Selection:** While gamma distribution is the standard choice for SPI calculation (used in our implementation), xclim supports alternative distributions including Pearson Type III and empirical distributions.
2. **Calculation Methods:** The library supports multiple calculation approaches: - "APP": Approximation method (used in our implementation) - "ML": Maximum likelihood estimation - "PWM":

Probability weighted moments

3. **Calibration Period Control:** As seen in the implementation, the function allows for specific calibration periods (`cal_start` and `cal_end`), crucial for establishing a consistent reference climatology across both forecast and observation datasets.
4. **Zero Precipitation Handling:** The `fitkwargs={"floc": 0}` parameter ensures that gamma distribution fitting properly handles zero precipitation values, which can be common in arid regions.

2.3 The xarray Foundation: Multi-dimensional Data Structures

Underlying both the data processing and SPI calculation is the xarray library, which provides labeled multi-dimensional arrays designed specifically for scientific datasets. In the context of drought monitoring, xarray offers several advantages:

1. **Self-describing Data:** Coordinates, dimensions, and metadata travel with the data, simplifying the tracking of spatial and temporal information.
2. **Advanced Indexing:** The implementation leverages xarray's powerful selection capabilities (e.g., `.sel(lead=fm)`, `.sel(latitude=slice(lat_max, lat_min))`) for intuitive data subsetting.
3. **Parallel Computing Integration:** Through dask integration, xarray enables scalable processing of large ensembles, evident in the `.compute()` calls throughout the processing pipeline.
4. **Interoperability:** The seamless handling of different data formats (NetCDF, GRIB, Zarr) facilitates the integration of observation data (CHIRPS) with forecast data (SEAS5).

In the SPI calculation for ensemble forecast data, the processing handles ensemble members differently based on their characteristics. For older members (0-24), SPI is calculated using the 1991-2018 calibration period, while newer operational members (25-51) use a 2017-2024 calibration period, addressing the challenge of maintaining consistent statistical properties across the full ensemble.

Regridding Datasets with xESMF

3.1 Spatial Resolution Harmonization

A critical step in comparing forecasts with observations involves addressing the inherent resolution mismatch between datasets. In the drought anticipatory action system, CHIRPS observations (native resolution of 0.05°) and SEAS5 forecasts (native resolution of 1.0°) must be harmonized to a common grid to enable meaningful comparisons. The `run-process-spi.py` script implements this harmonization through regridding, bringing both datasets to an intermediate resolution of 0.25° (approximately 25 km at the equator).

The regridding process handles two distinct cases:

1. CHIRPS Observation Regridding (5 km $\rightarrow$ 25 km):

```
# Create output grid with regular spacing
ds_out = xr.Dataset({
    "lat": ([ "lat" ], np.arange(lat_min, lat_max, res), {"units":
    "degrees_north"}),
    "lon": ([ "lon" ], np.arange(lon_min, lon_max, res), {"units":
    "degrees_east"}),
})

# Regrid using conservative method
regridder = xe.Regridder(ds1, ds_out, "conservative")
dr_out = regridder(dr, keep_attrs=True)
```

This upscaling process preserves precipitation totals through conservative regridding, which maintains the area-integrated rainfall quantity during the transition to a coarser grid.

2. SEAS5 Forecast Regridding (100 km $\rightarrow$ 25 km):

```
# Create output grid matching observations
ds_out = xr.Dataset({
    "lat": ([ "lat" ], kn_obs['lat'].values, {"units": "degrees_north"}),
    "lon": ([ "lon" ], kn_obs['lon'].values, {"units": "degrees_east"}),
})

# Perform regridding
regridder = xe.Regridder(gd2, ds_out, "bilinear", periodic=False)
dr_out = regridder(agd, keep_attrs=True)
```

For the forecast data, bilinear interpolation is employed to downscale from 100 km to 25 km, producing smoother transitions between grid cells appropriate for the forecast uncertainty.

3.2 The xESMF Library

The Earth System Modeling Framework (ESMF) regridding functionality is made accessible through xESMF, a Python interface that integrates seamlessly with xarray datasets. This library provides multiple interpolation methods appropriate for different variables and scientific needs:

1. **Conservative:** Preserves the integral of the field (used for CHIRPS precipitation upscaling)
2. **Bilinear:** Weighted average of nearest grid points (used for SEAS5 interpolation)
3. **Nearest-neighbor:** Assigns value from nearest source point
4. **Higher-order patch:** Smoother interpolation useful for continuous fields

In the implementation, regridding objects are created once and stored for repeated application across dimensions:

```
regridder = xe.Regridder(gd2, ds_out, "bilinear", periodic=False)
```

The `periodic=False` parameter indicates that the longitude dimension is not treated as periodic (wrapped around the globe), which is appropriate for regional domains in East Africa.

3.3 Handling Special Cases in Regridding

The regridding implementation addresses several technical challenges:

1. **Coordinate System Compatibility:** The script ensures coordinate naming consistency before regridding:

```
# Rename coordinates for consistency if needed
if 'longitude' in ds_p_m1.dims and 'latitude' in ds_p_m1.dims:
    gd2 = ds_p_m1.rename({'longitude': 'lon', 'latitude': 'lat'})
```

2. **Land-Sea Mask Preservation:** The regridding process preserves missing values to maintain the integrity of land-sea boundaries.
3. **Ensemble Dimension Handling:** When processing SEAS5 ensemble data, regridding is applied individually to each lead time to manage memory efficiently:

```
for fm in range(min(6, len(all_leads))):
    # Regridding for each lead time
    # ...
```

This approach ensures that both datasets align on a common 0.25° grid, enabling direct comparison between forecasts and observations while balancing spatial detail against computational efficiency in the SPI calculation and verification process.

Aligning Forecast and Observation Data with climpred

4.1 Time Alignment Challenge in Forecast Verification

One of the most complex challenges in verifying seasonal forecasts is aligning the time coordinates between forecasts and observations. The `make_obs_fct_dataset` function in `vthree_utils.py` addresses this critical task by integrating the SEAS5 forecast and CHIRPS observation datasets with different temporal structures:

```
def make_obs_fct_dataset(params):  
    """  
    Prepares observed and forecasted dataset subsets for a specific region,  
    season, and lead time by aligning their temporal coordinates.  
    """
```

The function performs a sophisticated alignment process that accounts for three distinct time dimensions in seasonal forecasting:

1. **Initialization Time (init):** When the forecast was issued
2. **Lead Time (lead):** How many months into the future the forecast extends
3. **Valid Time:** When the forecast is valid for (init + lead)

Since CHIRPS observations have only a `time` dimension while SEAS5 forecasts have both `init` and `forecastMonth` (lead time), direct comparison requires creating a compatible coordinate system.

4.2 The climpred Library

The alignment utilizes `climpred`, a specialized Python library for seasonal forecast verification that bridges the gap between forecast and observational data structures:

```
hindcast = HindcastEnsemble(a_fc)  
hindcast = hindcast.add_observations(a_obs)  
a_fcl = hindcast.get_initialized()
```

This sequence of operations performs several crucial tasks:

1. Creates a `HindcastEnsemble` object from the forecast dataset
2. Adds the observation dataset as a reference
3. Computes valid times for each forecast initialization and lead time combination
4. Realigns the forecast dataset with explicit `valid_time` coordinates

The `HindcastEnsemble` framework enables consistent handling of forecast lead times across multiple initialization dates, crucial for calculating skill metrics.

4.3 Lead Time Calculation Example: JJA Season SPI

To illustrate the lead time concept, consider forecasting the June-July-August (JJA) SPI:

Forecast Month	Lead Time	Valid Period (SPI JJA)	Practical Usage
March	Lead time 4	June-July-August	Early planning
April	Lead time 3	June-July-August	Refinement
May	Lead time 2	June-July-August	Final decision

For example, forecasts initialized in March (Lead time 4) need three months to reach the start of the JJA period, providing earlier warning but typically with lower skill. Conversely, May forecasts (Lead time 2) offer higher accuracy but less preparation time.

The alignment process must account for these varying lead times, extracting the correct forecast-observation pairs:

```
# Select forecast for specific lead time
a_fc2 = a_fc1.isel(lead=params.lead_int)

# Find common dates between forecasts and observations
common_dates = np.intersect1d(fct_valid_times, obs_times)

# Filter observations to include only dates with matching forecasts
a_obs3 = a_obs2.sel(time=common_dates)

# Calculate initial dates by subtracting lead time from valid dates
a_fc3_init_dates = common_dates.astype("datetime64[M]") - np.timedelta64(
    int(a_fc3.lead.values), "M"
)

# Select forecasts initialized on these dates
a_fc4 = a_fc3.sel(init=a_fc3_init_dates)
```

This process creates perfect temporal alignment between forecast valid times and observation times, enabling direct comparison for verification statistics. It ensures that when calculating metrics like hit rates or false alarm ratios, the system is comparing forecasts and observations for exactly the same temporal periods.

4.4 Seasonal SPI Product Identification

Another sophisticated aspect of the alignment is identifying specific SPI seasonal products from the time coordinates:

```
# Create season identifiers (e.g., "MAM", "JJA")
spi_prod_list = spi3_prod_name_creator(a_fc2, "valid_time")
obs_spi_prod_list = spi3_prod_name_creator(a_obs, "time")

# Assign season identifiers as coordinates
a_fc2 = a_fc2.assign_coords(spi_prod=("init", spi_prod_list))
a_obs1 = a_obs.assign_coords(spi_prod=("time", obs_spi_prod_list))

# Filter to the specified season
a_fc3 = a_fc2.where(a_fc2.spi_prod == params.season_str, drop=True)
a_obs2 = a_obs1.where(a_obs1.spi_prod == params.season_str, drop=True)
```

This approach enables targeting specific three-month periods like MAM (March-April-May) or JJA (June-July-August), aligning with key agricultural seasons in East Africa while maintaining the temporal correspondence between forecast and observation datasets.

Forecast Verification with xhistogram

5.1 Multi-dimensional Contingency Analysis

Forecast verification forms the cornerstone of any anticipatory action system, providing evidence that the forecasting system can reliably detect drought events before they occur. The `vthree_utils.py` script implements comprehensive verification methodology using the `xhistogram` library (imported as `xhist`), which enables efficient calculation of contingency tables and skill metrics across multiple dimensions.

At the core of this verification approach is the `xhist_metrics_2d` function, which performs grid-based verification of forecast probabilities against binary drought observations:

```
def xhist_metrics_2d(obs_data, ens_prob_data, params, calculate_auroc=True):
    """
    Calculates verification metrics for ensemble drought forecasts using
    multi-dimensional contingency tables.
    """
    # Get threshold for the specified region and season
    threshold_dict = get_threshold(params.region_id, params.sc_season_str)
    threshold = threshold_dict[params.level]

    # Process trigger values and calculate metrics
    df = process_trigger_values(
        obs_data, ens_prob_data, params, threshold, calculate_auroc
    )

    return df
```

This approach leverages multi-dimensional histograms to efficiently calculate contingency tables across the entire spatial domain, without requiring explicit loops over each grid cell.

5.2 Constructing Contingency Tables

The key innovation in this verification framework is the use of `xarray`'s histogram functionality to construct contingency tables directly. This is implemented in the `calculate_contingency_table` function:

```
def calculate_contingency_table(obs_event, forecast_event, params):
    """Creates a 2x2 contingency table using xhistogram"""
    obs_event1 = obs_event[params.spi_prod_name]
    forecast_event1 = forecast_event[params.spi_prod_name]

    # Align forecast to observation time
```

```

forecast_event3 = forecast_event1.assign_coords(
    init=forecast_event1.coords["valid_time"]
)
forecast_event3 = forecast_event3.rename({"init": "time"})

# Histogram calculation produces the contingency table
contingency_table = xhist.histogram(
    obs_event1, forecast_event3, bins=[2, 2], density=False, dim=["lat",
"lon"]
)

return contingency_table

```

This function:

1. Extracts the relevant SPI values from the observation and forecast datasets
2. Aligns the temporal coordinates of forecasts with observations
3. Computes a 2D histogram where each axis represents binary values (0=no drought, 1=drought)
4. Returns a contingency table with dimensions [time, obs_event, forecast_event]

The resulting contingency table has four elements representing the classic “hits”, “misses”, “false alarms”, and “correct negatives” for each time step.

5.3 Verification Metrics Explained

From the contingency tables, a comprehensive set of verification metrics is calculated to assess different aspects of forecast quality:

1. Hit Rate (HR):

$$HR = \frac{\text{hits}}{\text{hits} + \text{misses}}$$

Measures the proportion of observed drought events that were correctly forecast. Also called Probability of Detection (POD). Ranges from 0 (worst) to 1 (perfect).

2. False Alarm Ratio (FAR):

$$FAR = \frac{\text{false_alarms}}{\text{false_alarms} + \text{hits}}$$

Measures the proportion of drought warnings that were false alarms. Ranges from 0 (perfect) to 1 (worst).

3. Bias Score (BS):

$$BS = \frac{\text{hits} + \text{false_alarms}}{\text{hits} + \text{misses}}$$

Indicates whether the system tends to under-forecast ($BS < 1$) or over-forecast ($BS > 1$) drought events.

4. Hanssen-Kuipers Score (HK):

$$HK = HR - \frac{\text{false_alarms}}{\text{false_alarms} + \text{correct_negatives}}$$

Also known as the True Skill Statistic (TSS), measures the ability to discriminate between drought and non-drought events. Ranges from -1 to 1.

5. Heidke Skill Score (HSS):

$$HSS = \frac{2(\text{hits} \cdot \text{correct_negatives} - \text{misses} \cdot \text{false_alarms})}{(\text{hits} + \text{misses})(\text{misses} + \text{correct_negatives}) + (\text{hits} + \text{false_alarms})(\text{false_alarms} + \text{correct_negatives})}$$

Measures the improvement of the forecast over random chance. Ranges from $-\infty$ to 1.

6. Area Under ROC Curve (AUROC):

Measures the forecast's ability to discriminate between drought and non-drought events across all possible threshold values. Ranges from 0 to 1, with 0.5 indicating no skill.

5.4 AUROC and Bootstrapping

The AUROC calculation with bootstrapping represents the most sophisticated element of the verification methodology:

```
def calculate_auroc_score(contingency_table, trigger_value, n_bootstrap=1000):
    """
    Calculates the Area Under the ROC curve with confidence intervals
    using bootstrapping.
    """
    # Sum the contingency table across time
    summed_table = contingency_table.sum(dim="time")

    # Extract counts
    ct = summed_table.values.flatten()
    correct_negatives, false_alarms, misses, hits = ct
    total = hits + false_alarms + misses + correct_negatives

    # Perform bootstrapping
    auroc_bootstrap_scores = []
    for _ in range(n_bootstrap):
        bootstrap_counts = np.random.multinomial(
            total,
            [
                hits / total,
                misses / total,
                false_alarms / total,
                correct_negatives / total,
            ],
            size=1,
        )

        bootstrap_hits, bootstrap_misses, bootstrap_false_alarms,
        bootstrap_correct_negatives = bootstrap_counts[0]

        auroc_bootstrap_scores.append(
            calculate_auroc(
                bootstrap_hits,
                bootstrap_misses,
                bootstrap_false_alarms,
                bootstrap_correct_negatives,
            )
        )

    # Calculate mean and confidence intervals
    auroc_score = np.mean(auroc_bootstrap_scores)
    auroc_lb, auroc_ub = np.percentile(auroc_bootstrap_scores, [2.5, 97.5])

    return {
        "trigger_value": trigger_value,
        "auroc_score": auroc_score,
        "auroc_lb": auroc_lb,
        "auroc_ub": auroc_ub,
    }
```

This bootstrap approach:

1. Resamples from the original contingency table 1000 times
2. Calculates the AUROC score for each resample
3. Computes the mean AUROC score and 95% confidence intervals

Bootstrapping is crucial for drought verification because:

- Droughts are relatively rare events, leading to sampling uncertainty
- Confidence intervals provide stakeholders with a measure of uncertainty
- It helps identify whether skill scores are statistically significant

5.5 2D vs 1D Verification Methods

The system implements two complementary verification approaches:

1. 2D Verification (Grid-based):

```
def run_xhist2d(params):  
    """Grid-based verification across the entire spatial domain"""  
    # ...  
    df = xhist_metrics_2d(obs_data, fct_mod, params, calculate_auroc=True)  
    # ...
```

This approach calculates metrics at each grid point and then aggregates the results, preserving the spatial structure of forecast skill. It's useful for understanding regional variations in forecast quality.

2. 1D Verification (Region-based):

```
def run_xhist1d(params):  
    """Region-based verification using spatial means"""  
    # ...  
    obs_df = mean_obs_spi(obs_data, params.spi_prod_name)  
    fct_df = mean_emp_prob(fct_mod, fct_sev, fct_ext, params.spi_prod_name)  
    # ...
```

This approach first calculates spatial means across the region and then performs verification on these mean values. It's useful for making operational decisions at the regional level.

Both approaches are necessary because:

- 2D verification captures the spatial heterogeneity of forecast skill
- 1D verification aligns better with decision-making at administrative levels
- Comparing both can identify issues with compensating errors in the forecasting system

The comprehensive verification framework implemented in these scripts provides the statistical basis for selecting optimal trigger values for anticipatory action, balancing the desire for high hit rates against the need to minimize false alarms within operational constraints.

Trigger Selection for Anticipatory Action against Droughts in East Africa

6.1 Forecast Verification and Trigger Selection Process

6.1.1 Data Processing Pipeline

1. Initial Processing

- CHIRPS rainfall observations are processed to calculate SPI-3
- SEAS5 precipitation forecasts are processed to calculate SPI-3 for different lead times
- Both datasets are spatially aligned and masked to the region of interest (e.g., Karamoja)

2. Verification Workflow

- `run_xhist2d.py` performs 2D verification by comparing ensemble SPI forecasts against observed SPI
- `run_xhist1d.py` performs 1D verification by comparing mean ensemble SPI against regional mean SPI
- Results are saved as CSV files for different lead times, regions, and seasons

6.1.2 Critical Manual Step: Trigger Selection

The most critical manual step in the workflow is **trigger selection**, which occurs after verification metrics are calculated. This process:

1. Examines the CSV files generated by `run_xhist2d.py` and `run_xhist1d.py`
2. Identifies optimal trigger values based on verification metrics
3. Manually marks selected triggers by adding a value of '1' in the `td` column

For example, in the CSV file for Karamoja MAM season with lead time 3:

```
tv year# hr far bs hk hs au odc h m FA CN h% cat td
15.2 10 96.9 1.6 1.0 -0.0 -0.0 1.0 13 12 1 27 2 92.3 mild 0.0
...
68.0 21 72.7 0.0 0.7 nan 0.0 0.7 8 3 5 17 17 37.5 mod 1.0
...
```

The row with `td=1.0` indicates the manually selected trigger value (68% probability) for the moderate category.

6.1.3 Selection Criteria

Triggers are selected based on a combination of metrics:

- **AUROC** (Area Under ROC Curve) > 0.5 (skill better than random)
- **Hit Rate** (HR) > 50% (minimum detection rate)
- **False Alarm Ratio** (FAR) < 35% (maximum acceptable false alarms)
- **Operational considerations** (lead time, practical implementation)

6.2 Verification Metrics

6.2.1 2D Verification Metrics

- Used for spatial analysis of forecast skill
- Calculates contingency table statistics by gridcell
- Produces `auroc_scores`, `hit_rates`, and `false_alarm_ratios`

6.2.2 1D Verification Metrics

- Analyzes regional mean forecast performance
- Used for operational decision-making
- Focuses on practical application of triggers
- Metrics include:
 - `hit_rate`: Percent of actual drought events correctly forecast
 - `false_alarm_ratio`: Proportion of forecast events that didn't occur
 - `bias_score`: Ratio of forecast frequency to observed frequency
 - `hanssen_kuipers_scores`: Difference between hit rate and false alarm rate
 - `heidke_skill_scores`: Forecast accuracy relative to random chance

6.3 Downstream Processing Dependencies

The manually selected triggers (marked with `td=1`) are essential inputs for:

1. `run_bar_plot.py`

- Creates comparative bar visualizations
- Shows observed SPI values alongside forecast probabilities
- Highlights performance based on selected trigger values

2. `run_heatmap.py`

- Generates heatmaps of hit rates and false alarm ratios
- Visualizes performance across different months and lead times
- Supports seasonal decision-making

3. `run_latex_table.py`

- Produces formatted tables of available triggers
- Highlights selected triggers
- Documents full decision metrics for stakeholders

6.4 Forecast Skill Assessment

For each potential trigger, the system calculates:

- Number of hits (h): Correctly forecast drought events
- Number of misses (m): Unforecast drought events
- Number of false alarms (FA): Forecast events that didn't occur
- Number of correct negatives (CN): Correctly forecast non-events
- Hit percentage (h%): Percentage of observed events successfully forecast

6.5 Operational Implementation

The selected triggers enable a three-phase anticipatory action approach:

1. **Ready:** Initial alert based on early forecasts
2. **Set:** Confirmation with subsequent forecasts
3. **Go:** Implementation of anticipatory actions

The triggers serve as decision thresholds for initiating the anticipatory action workflow, with different trigger values potentially applied based on vulnerability levels (general vs. emergency menus).

6.6 Time Window Considerations

Triggers are optimized for specific seasonal windows:

- **Window 1:** Start to mid-season (e.g., MAM for Karamoja)
- **Window 2:** Middle to end of season (e.g., JJA for Karamoja)

This temporal approach allows for targeted interventions at critical points in the agricultural cycle.

Map Visualization for Ensemble Forecasts and Observations

This document explains how to create map visualizations from ensemble forecasts and observational data for drought monitoring using Python GIS tools.

7.1 Setting Up the Environment

First, import the necessary libraries:

```
import xarray as xr
import numpy as np
import matplotlib.pyplot as plt
import cartopy.crs as ccrs
import geopandas as gpd
import os
from PIL import Image
from shapely import wkb
```

7.2 Data Organization with DataTree

DataTree provides a hierarchical structure for organizing ensemble data efficiently:

```
def helper_stamp_plot(ens_data, obs_data, fct_mod, fct_sev, fct_ext):
    """
    Helper function to organize data into a tree structure.

    Parameters:
    ens_data (xarray.Dataset): Ensemble data
    obs_data (xarray.Dataset): Observation data
    fct_mod, fct_sev, fct_ext (xarray.Dataset): Forecast datasets
    """
    seas51tree = xr.DataTree()

    # Add ensemble members
    for member in ens_data.member:
        member_data = ens_data.sel(member=member)
        seas51tree[f"ensemble/member_{int(member)}"] = xr.DataTree(
            name=f"member_{int(member)}",
            dataset=member_data
        )
```

Add observation and forecast datasets to the DataTree:

```
# Add observation and forecast datasets
```

```
seas5ltree["observation"] = xr.DataTree(
    name="observation",
    dataset=obs_data
)
seas5ltree["fct_mod"] = xr.DataTree(
    name="fct_mod",
    dataset=fct_mod
)
seas5ltree["fct_sev"] = xr.DataTree(
    name="fct_sev",
    dataset=fct_sev
)
seas5ltree["fct_ext"] = xr.DataTree(
    name="fct_ext",
    dataset=fct_ext
)

return seas5ltree
```

7.3 Creating Single-Row Time Series Plot

The following function creates visualizations for a specific initialization date:

```
def create_single_row_plot(tree, init, params, region_geom):
    """
    Create a row of maps for a specific initialization date.

    Parameters:
    tree (xarray.DataTree): Data organized in tree structure
    init (datetime): Initialization date
    params (object): Parameters for plotting
    region_geom (geometry): Region geometry for overlay
    """
    members = list(tree["ensemble"].children.keys())
    valid_times = tree["ensemble/member_0"].ds.valid_time.values
    num_members = len(members)
    num_additional_plots = 4 # Obs, mod, sev, ext
    total_plots = num_members + num_additional_plots
```

Set up the figure and define color scales:

```
# Create figure with map projections
fig, axs = plt.subplots(
    1, total_plots,
    figsize=(2 * total_plots, 2),
    subplot_kw={"projection": ccrs.PlateCarree()})

# Define color scales
ensemble_cmap_range = (-4, 4)
fct_cmap_range = (0.0, 1.0)

# Get valid time for this initialization
valid_time = valid_times[
    np.where(tree["ensemble/member_0"].ds.init.values == init)[0][0]
]
```

Plot ensemble members and additional data:

```
# Plot ensemble members
```

```

for j, member_key in enumerate(members):
    _plot_ensemble_member(
        tree, member_key, params.spi_prod_name, init,
        axs[j], ensemble_cmap_range, region_geom,
        shape_overlay=(j == 0)  # Only show shapefile on first plot
    )

# Add observation and forecast plots
plot_titles = ["Obs", "mod", "sev", "ext"]
plot_keys = ["observation", "fct_mod", "fct_sev", "fct_ext"]

for k, (title, key) in enumerate(zip(plot_titles, plot_keys)):
    _plot_additional_data(
        tree, key, params.spi_prod_name, init, valid_time,
        axs[num_members + k], title,
        ensemble_cmap_range if key == "observation" else fct_cmap_range
    )

return fig, axs, (2 * total_plots, 2)

```

7.4 Plotting Map Elements

The system uses helper functions to plot ensemble members:

```

def _plot_ensemble_member(tree, member_key, variable, init, ax, cmap_range,
                          geom, shape_overlay=False):
    """Plot individual ensemble member."""
    member_data = tree[f"ensemble/{member_key}"].ds[variable]
    data = member_data.sel(init=init).values
    time_np64 = np.array(str(init), dtype="datetime64[ns]")
    year = pd.to_datetime(time_np64).year

    # Plot data
    ax.pcolormesh(
        tree["ensemble/member_0"].ds.lon.values,
        tree["ensemble/member_0"].ds.lat.values,
        data, cmap="RdBu", transform=ccrs.PlateCarree(),
        vmin=cmap_range[0], vmax=cmap_range[1]
    )

```

Add title, formatting, and optional shapefile overlay:

```

# Add title and formatting
ax.set_title(f'{year}m{member_key.split("_")[1]}', fontsize=6)
ax.set_xticks([])
ax.set_yticks([])

# Overlay shapefile if requested
if shape_overlay:
    ax.add_geometries(
        [geom], crs=ccrs.PlateCarree(),
        edgecolor="black", facecolor="none"
    )

```

Helper function to plot observation or forecast data:

```

def _plot_additional_data(tree, key, variable, init, valid_time, ax,
                          title, cmap_range):
    """Plot observation or forecast data."""
    dataset = tree[key].ds[variable]

```

```

# Handle different coordinate naming conventions
coord_key = "time" if "time" in dataset.coords else "init"
obs_init = (
    np.datetime64(valid_time.strftime("%Y-%m-%d %H:%M:%S"))
    if coord_key == "time" else init
)

# Extract data
obs_data = dataset.sel({coord_key: obs_init}).values

```

Set appropriate colormap and plot data:

```

# Set appropriate colormap
cmap = "RdBu" if key == "observation" else "Blues"

# Plot data
ax.pcolormesh(
    tree["ensemble/member_0"].ds.lon.values,
    tree["ensemble/member_0"].ds.lat.values,
    obs_data, cmap=cmap, transform=ccrs.PlateCarree(),
    vmin=cmap_range[0], vmax=cmap_range[1]
)

# Add title and formatting
ax.set_title(title, fontsize=6)
ax.set_xticks([])
ax.set_yticks([])

```

7.5 Processing Multiple Initializations

To create a complete visualization, the system processes multiple initialization dates:

```

def plot_allrows(seas5ltree, shapefile_df, params):
    """
    Create plots for all initialization dates and combine them.

    Parameters:
    seas5ltree (xarray.DataTree): Data in tree structure
    shapefile_df (GeoDataFrame): Shapefile data
    params (object): Plot parameters
    """
    # Create output directory
    output_dir = (f"{params.output_path}map_{params.region_id}_"
                  f"{params.sc_season_str}_lt{params.lead_int}")
    os.makedirs(output_dir, exist_ok=True)

    # Get region geometry
    region_geom = shapefile_df["geometry"].values[0]

    # Get all initialization dates
    inits = seas5ltree["ensemble/member_0"].ds.init.values
    plot_files = []

```

Create plots for each initialization date:

```

# Process each initialization date
for i, init in enumerate(inits):
    fig, axs, figsize = create_single_row_plot(
        seas5ltree, init, params, region_geom
    )

```

```

    )

    # Save individual plot
    output_file = f'{output_dir}/stamp_plot_{init.strftime("%Y%m%d")}.png'
    plot_files.append(output_file)
    fig.savefig(output_file, dpi=100, bbox_inches="tight")
    plt.close(fig)

    # Add colorbar and title to the final figure
    add_colorbar_and_title(figsize, params, output_dir)

    # Merge all plots into a single image
    merge_png_files(
        input_dir=output_dir,
        output_file=(f"map_{params.region_id}_{params.sc_season_str}_"
                     f"lt{params.lead_int}.pdf"),
        delete_originals=False
    )

```

7.6 Merging Individual Plots

The following function combines individual timestep plots into a single visualization:

```

def merge_png_files(input_dir, output_file, delete_originals=False):
    """
    Merge individual PNG files into a single image.

    Parameters:
    input_dir (str): Directory containing PNG files
    output_file (str): Name of output file
    delete_originals (bool): Whether to delete original files
    """
    # Get all PNG files
    png_files = sorted(Path(input_dir).glob("*.png"))

    # Open all images and get dimensions
    images = []
    max_width = 0
    total_height = 0

```

Process images and create a new composite:

```

for png_file in png_files:
    with Image.open(png_file) as img:
        images.append(img.copy())
        max_width = max(max_width, img.width)
        total_height += img.height

    # Create new composite image
    merged_image = Image.new("RGB", (max_width, total_height), (255, 255, 255))

    # Paste each image into the merged image
    y_offset = 0
    for img in images:
        merged_image.paste(img, ((max_width - img.width) // 2, y_offset))
        y_offset += img.height

    # Save merged image
    output_path = os.path.join(input_dir, output_file)
    merged_image.save(output_path, dpi=(300, 300))

```

Optionally delete original files:

```
# Optionally delete original files
if delete_originals:
    for png_file in png_files:
        os.remove(png_file)
```

7.7 Example Usage

To run the map visualization process:

```
def run_map_plot(params, use_local=False):
    """
    Run the complete mapping process.

    Parameters:
    params (object): Parameters for the run
    use_local (bool): Whether to use local files

    Returns:
    str: Status message
    """
    # Get thresholds for the region and season
    threshold_dict = get_threshold(params.region_id, params.sc_season_str)

    # Prepare datasets
    obs_data, ens_data = make_obs_fct_dataset(params)
```

Process data and create plots:

```
# Calculate empirical probabilities
fct_mod, fct_sev, fct_ext = seas5l_patch_empirical_probability(
    ens_data, threshold_dict
)

# Create data tree
dstree = helper_stamp_plot(ens_data, obs_data, fct_mod, fct_sev, fct_ext)

# Get region bounds
if use_local:
    shapefile_df, extent = get_region_bounds(region_id, use_local=True)
else:
    credentials = get_credentials('coiled-data.json')
    shapefile_df, extent = get_region_bounds(
        region_id, credentials, use_local=False, buffer=0.5
    )

# Create and merge plots
plot_allrows(dstree, shapefile_df, params)
```

Create final merged PDF:

```
# Create final merged PDF
merge_png_files(
    input_dir=(f"{params.output_path}map_{params.region_id}_"
               f"{params.sc_season_str}_lt{params.lead_int}"),
    output_file=(f"map_{params.region_id}_{params.sc_season_str}_"
                 f"lt{params.lead_int}.pdf"),
    delete_originals=False
)
```

```
return "merge plot"
```

7.8 Conclusion

This approach allows for the effective visualization of both ensemble forecasts and observed data on maps with region overlays. The plots are organized by initialization date, allowing stakeholders to track forecast evolution over time and make informed decisions about anticipatory actions for drought mitigation.

